


Cow in a Box & Friends

a tale of smart pointers

Rust Wrocław



About me

- Magister inżynier
- Software developer



Agenda

- Pointer
- Smart Pointer
- Container
- Interior Mutability
- ~~Multithreading~~


Pointer

- Rust does have raw pointers...
 - both **const** and **mut**

```
fn main() {  
    let x = 6;  
    let px: *const i32 = &x;  
    println!("{:?}", px);  
}
```

```
magister@inzynier:~/rust$ cargo run
```

```
0x7fff5277361c
```




Pointer

- ...but they cannot be **safely** dereferenced

```
fn main() {  
    let x = 6;  
    let px: *const i32 = &x;  
    println!("{}", *px);  
}
```

```
magister@inzynier:~/rust$ cargo run
```

```
|     println!("{}", *px);  
|                               ^^^ dereference of raw pointer
```



Pointer


- **Unsafe** to the rescue!

```
let x = 6;  
let px: *const i32 = &x;  
unsafe {  
    println!("{}", *px);  
}  
// -or just-  
println!("{}", unsafe { *px });
```

```
magister@inzynier:~/rust$ cargo run
```

```
6
```

```
6
```



Smart Pointer

- It's all about ownership
- Who owns a piece of data?
- Owner is obliged to release the memory:
 - safely
 - not prematurely
 - at all



Smart Pointers

- Box $(\sim T)$
- Rc $(@T)$
- Arc
- Weak
- Cow
- ...*some more?*



Smart Pointers - Box

- Keeps exclusive ownership of the memory
- If you can't or doesn't want to keep a value on the stack
 - Its size is not known at compile time (like trait objects, for example)
 - It's too big
- Implements **Deref** trait
 - It shines in conjunction with **deref coercion**
- Mutations allowed

```
fn main() {  
    let one = 1;  
    let two = Box::new(2);  
  
    let r_one = &one;  
    println!("{}", *r_one, *two);  
}
```

```
let a = box Point  
{  
    x: 10, y: 20  
};
```

← Former form :)

```
magister@inzynier:~/rust$ cargo run
```

```
1 - 2
```

Smart Pointers - Rc

- Shares the ownership of the memory
- Data must remain valid as long as there is at least one owner
- To be used in single-threaded cases only
 - The internal reference counter **is not** atomic
- Mutations **not** allowed (no single, clear owner)

```
let mut one = Box::new(1);  
*one = 2;  
println!("{}", *one);
```

```
~/rust$ cargo run
```

```
2
```

```
let mut one = Rc::new(1);  
*one = 2;  
println!("{}", *one);
```

```
~/rust$ cargo run
```

```
| *one = 2;  
| ^^^^^^^^ cannot assign
```

Smart Pointers - Rc

```
struct Data {  
    small: i32,  
    huge: Rc<String>,  
}  
  
impl Data {  
    pub fn new() -> Self { Data { small: 1,  
        huge: Rc::new("Loooong".to_string(),),}}  
  
    pub fn get_small(&self) -> i32 { self.small }  
    pub fn get_huge(&self) -> Rc<String> {  
        Rc::clone(&self.huge)}  
}
```

```
fn main() {  
    let db = Data::new();  
    let (x, y, z) = (db.get_small(), db.get_small(),  
        db.get_small());  
    println!("{}", x, y, z);  
    let (x, y, z) = (db.get_huge(), db.get_huge(),  
        db.get_huge());  
    println!("{}", x, y, z);  
    println!("{}", Rc::strong_count(&y),  
        Rc::ptr_eq(&x, &z));  
    std::mem::drop(db);  
    println!("{}", Rc::strong_count(&y));  
}
```

```
magister@inzynier:~/rust$ cargo run  
1 1 1  
Loooong Loooong Loooong  
4 - true  
3
```

Smart Pointers - Rc

- Usage across threads

```
struct Data {}

fn main() {
    let data = Rc::new(Data {});
    thread::spawn(move || {
        let local = Rc::clone(&data);
    });
}
```

```
magister@inzynier:~/rust$ cargo build
```

```
| `std::rc::Rc<Data>` cannot be sent between threads safely
```

Smart Pointers - Arc

- Usage across threads
- Now we're safe, but we need to explicitly indicate that we are ready to pay the cost of synchronization
- Contained **T** must be **Send** and **Sync**

```
struct Data {}  
  
fn main() {  
    let data = Arc::new(Data {});  
    thread::spawn(move || {  
        let local = Arc::clone(&data);  
    });  
}
```

```
magister@inzynier:~/rust$ cargo build
```

```
Finished dev [unoptimized + debuginfo] target(s) in 0.02s
```

Smart Pointers - Weak

- `std::rc::Weak` and `std::sync::Weak`
 - Result of downgrading **Rc** and **Arc** respectively
- They are smart in a way that they can “tell” if reference is still valid

```
struct Data {}  
fn main() {  
    let weak;  
    {  
        let data = Rc::new(Data {});  
        weak = Rc::downgrade(&data);  
    }  
    let data = weak.upgrade();  
    assert!(data.is_none());  
}
```

```
struct Data {}  
fn main() {  
    let data = Rc::new(Data {});  
    let weak;  
    {  
        weak = Rc::downgrade(&data);  
    }  
    let data = weak.upgrade();  
    assert!(data.is_some());  
}
```

Smart Pointers - Cow

- Clone-on-Write
- Provides immutable access to data
- Create clone of referenced data when **mutability** or **ownership** is required
- Imagine accessing a text file from the read-only network share
 - Anyone can seamlessly open it and read the text
 - But if someone wishes to make a correction, he needs to create a local copy
- The same principles apply to data stored in **Cow**
 - There's just a single instance of the (potentially huge) data
 - Until someone wants to mutate it (take ownership)
 - Then the clone is made

Smart Pointers - Cow

```
fn main() {  
    let data = HugeData { number: 1 };  
  
    let mut ptr1 = Cow::Borrowed(&data);  
    println!("{}", ptr1.number);  
  
    let mut ptr2 = ptr1.to_mut();  
    ptr2.number = 2;  
    println!("{}", ptr2.number);  
}
```

```
impl Clone for HugeData {  
    fn clone(&self) -> Self {  
        println!("Well, here's your copy!");  
        HugeData {  
            number: self.number,  
        }  
    }  
}
```

```
magister@inzynier:~/rust$ cargo run
```

```
1
```

```
Well, here's your copy!
```

```
2
```


Smart Pointers - Cow

```
fn foo() -> &'static str {  
    let error_code = 17;  
    match error_code {  
        1 => "Not enough memory".into(),  
        2 => "Too much memory".into(),  
        _ => "Unknown".into(),  
    }  
}
```

- Returning static “strings” to avoid allocations
- Works great!

Smart Pointers - Cow

```
fn foo() -> &'static str {  
    let error_code = 17;  
    match error_code {  
        1 => "Not enough memory".into(),  
        2 => "Too much memory".into(),  
        _ => format!("Unknown: {error_code}").into(),  
    }  
}
```

- Works as long as we don't need to pass any additional details

```
magister@inzynier:~/rust$ cargo build  
8 |         _ => format!("Unknown: {}", error_code).into(),  
  |                                     ^^^^ the trait  
  | `std::convert::From<std::string::String>` is not implemented for `&str`
```

Smart Pointers - Cow

```
fn foo() -> String {  
    let error_code = 17;  
    match error_code {  
        1 => "Not enough memory".into(),  
        2 => "Too much memory".into(),  
        _ => format!("Unknown:  
{}", error_code).into(),  
    }  
}
```

- So, just return an owned string
- Works, but not optimal performance-wise

```
magister@inzynier:~/rust$ cargo build  
Finished dev [unoptimized + debuginfo] target(s) in 0.01s
```

Smart Pointers - Cow

```
fn foo() -> Cow<'static, str> {  
    let error_code = 17;  
    match error_code {  
        1 => "Not enough memory".into(),  
        2 => "Too much memory".into(),  
        _ => format!("Unknown:  
{}", error_code).into(),  
    }  
}
```

- Now, owned string is returned only when really needed

```
magister@inzynier:~/rust$ cargo build  
Finished dev [unoptimized + debuginfo] target(s) in 0.01s
```

Smart Pointers - Cow

```
impl<'a> From<&'a str> for Cow<'a, str> {  
    #[inline]  
    fn from(s: &'a str) -> Cow<'a, str> {  
        Cow::Borrowed(s)  
    }  
}
```

```
impl<'a> From<String> for Cow<'a, str> {  
    #[inline]  
    fn from(s: String) -> Cow<'a, str> {  
        Cow::Owned(s)  
    }  
}
```

- Cow will automatically pass on borrowed reference, until a Clone is needed
- Drop-in replacement in cases like that

Smart Pointers - Deref Coercion

- *“This rule is one of the only places in which Rust does an automatic conversion for you”* (for convenience)

If type **U** can be dereferenced to other type **T** then **&U** will automatically act as **&T**

If **Box<i32>** implements **Deref<Target=i32>** then **&Box** can be used in every place when **&i32** is expected.

&&&&&&&Box fits as well :-)

Smart Pointers - Deref Coercion

```
fn takes_ref(r: &i32) {  
    println!("{}", r);  
}  
  
fn main() {  
    let a = 7;  
    takes_ref(&a);  
  
    let b = Box::new(7);  
    takes_ref(&b);  
    takes_ref(&&&&&&&&&&&b);  
  
    takes_ref(&>(*b)); // Manual deref  
}
```

```
magister@inzynier:~/rust$ cargo run  
7  
7  
7  
7
```

Smart Pointers - Summary

- Choose your guarantees

	Box	Rc	Arc	Cow
Ownership	Single	Shared	Shared	Mixed
Mutability	✓	✗	✗	
Thread-safety	✓	✗	✓	

Inherited Mutability

- By design, Rust exhibits the **inherited** mutability
- Mutability is defined at binding-level (all or nothing approach)

```
struct S {  
    a: i32,  
    b: i32,  
}  
  
fn main() {  
    let s = S { a: 1, b: 2 };  
    let mut ms = S { a: 1, b: 2 };  
}
```

```
struct S {  
    a: mut i32,  
    b: const i32,  
}
```



Interior Mutability

- It might be required to mutate some members (mind the **mutable** keyword from C++)
- API should remain non-mutable
 - Do not force user to create mutable bindings just for the sake of some internal data manipulation
- Examples
 - Caching (Memoization)
 - Lazy-evaluation
 - Collecting performance statistics
 - Debugging
 - Increasing/decreasing reference counter in **Rc**

Interior Mutability - Toolbox

- UnsafeCell
 - The only core language feature to work around the mutability restrictions
 - Rust uses this basic tool as a foundation for more sturdy wrappers
 - A tool not to be used directly (!)
- Cell
 - A mutable memory location
 - Provides functionality corresponding to **mutable** class members in C++
 - Owns the memory location
- RefCell
 - Similar to Cell, but works on borrowed data
 - Provides access to **&mut** through **&self**
 - Never Sync

Interior Mutability - Cell

```
use std::cell::Cell;

struct S {
    a: i32,
    b: Cell<i32>,
}

fn main() {
    let s = S {
        a: 1,
        b: Cell::new(2),
    };
    s.b.set(3);
    println!("{}", s.b.get());
}
```

```
magister@inzynier:~/rust$ cargo run
3
```

Interior Mutability - Cell

- There is no borrow checking!

```
fn calculate_bonuses(hr_data: &Cell<i32>) {  
    let salary = hr_data.get();  
    hr_data.set(salary + 700);  
}  
  
fn main() {  
    let hr_database = Cell::new(15_000);  
    let salary = hr_database.get();  
    calculate_bonuses(&hr_database);  
    hr_database.set(salary);  
    println!("{}", hr_database.get());  
}
```

Prefer regular references or RefCell, so that Rust can protect you against such bugs.

```
~/rust$ cargo run  
15000
```

Interior Mutability - RefCell

```
use std::cell::RefCell;

struct S {
    b: RefCell<i32>,
}

fn main() {
    let s = S { b: RefCell::new(1) };
    let mut b = s.b.borrow_mut();
    *b = 2;
    println!("{}", b);
}
```


```
magister@inzynier:~/rust$ cargo run
```

```
2
```

“Fixing” the issues with borrow checker :)

```
fn main() {  
    let mut data = 23;  
    let r = &data;  
    println!("{}", r);  
    if get_user_input() > 10 {  
        let mr = &mut data;  
        *mr = 24;  
        println!("{}", mr);  
    }  
}
```

```
fn main() {  
    let mut data = 23;  
    let r = &data;  
  
    if get_user_input() > 10 {  
        let mr = &mut data;  
        *mr = 24;  
        println!("{}", mr);  
    }  
    println!("{}", r);  
}
```



“Fixing” the issues with borrow checker :)

```
error[E0502]: cannot borrow `data` as mutable because it is also borrowed as immutable
--> src/main.rs:9:18
|
7 |     let r = &data;
|           ----- immutable borrow occurs here
8 |     if get_user_input() > 10 {
9 |         let mr = &mut data;
|                   ^^^^^^^^^ mutable borrow occurs here
...
13 |     println!("{}", r);
|                   - immutable borrow later used here

error: aborting due to previous error
```


The “unsafe” hammer

```
fn main() {  
    unsafe {  
        let mut data = 23;  
        let r = &data;  
        if get_user_input() > 10 {  
            let mr = &mut data;  
            *mr = 24;  
            println!("{}", mr);  
        }  
        println!("{}", r);  
    }  
}
```



The “RefCell” hammer

```
fn main() {  
    let data = RefCell::new(23);  
    let r = data.borrow();  
    if get_user_input() > 10 {  
        let mut mr = data.borrow_mut();  
        *mr = 24;  
        println!("{}", mr);  
    }  
    println!("{}", r);  
}
```

Fixed ✓

```
magister@inzynier:~/rust$ cargo run  
23
```

The “RefCell” hammer

- It's fixed, until the user input remains 10 or less...
- Otherwise...

```
~/rust$ cargo run
    Running `target/debug/rust`
thread 'main' panicked at 'already borrowed: BorrowMutError',
/rustc/b8cedc00407a4c56a3bda1ed605c6fc166655447/src/libcore/cell.rs:878:9
note: run with `RUST_BACKTRACE=1` environment variable to display a backtrace
```

- **try_borrow()** family of functions can help

The “Rc<RefCell>” combo

- Rc - shared ownership
- RefCell - mutability...
- 🤔
- Not suitable in multi-threaded environment
 - **Arc**, **RwLock** or **Mutex**

```
fn main() {  
    let data = Rc::new(RefCell::new(23));  
    let r = data.clone();  
    if get_user_input() > 10 {  
        let mr = data.clone();  
        let mut internal = mr.borrow_mut();  
        *internal = 24;  
    }  
    println!("{}", *r.borrow());  
}
```

```
magister@inzynier:~/rust$ cargo run  
24
```

Summary

- Prefer inherited mutability (“cellless” code)
- Checklist before using interior mutability
 - Make sure you understand what you’re doing
 - **panic!** in runtime, even though safe, is still worse than compilation error
 - Rethink you approach, take 1, 2, ... 10 steps back

THE END